

PROFESSIONAL EXPERIENCE

Postdoctoral Associate October 2024 – Present
Center for Data Science, New York University, *Advised by He He*

Graduate Research Assistant September 2017 – September 2024
University of Massachusetts Amherst, *Advised by Andrew McCallum*

Research Intern June 2021 – September 2021
Google Research, *Mentored by Sandeep Tata*

Visiting Researcher June 2020 – August 2020
Chan Zuckerberg Initiative, *Mentored by Sunil Mohan*

Research Intern May 2017 – August 2017
MIT Lincoln Laboratory, *Mentored by Kyle O’Brien and Michael Yee*

Undergraduate Research Assistant May 2015 – May 2017
University of Michigan, *Advised by Grant Schoenebeck*

Undergraduate Research Assistant September 2013 – September 2015
University of Michigan, *Advised by Andrew DeOrio*

EDUCATION

University of Massachusetts, Amherst, MA
Ph.D., Computer Science, September 2024
Advisor: Andrew McCallum
Committee: Zico Kolter, Daniel Sheldon, Cameron Musco
Thesis: *Fast, Scalable, Warm-start Semidefinite Programming with Application to Knowledge Editing and Mixed-Integer Semidefinite Optimization*

University of Michigan, Ann Arbor, MI
B.S.E. Computer Science and Engineering, Minor in Mathematics, May 2017

PUBLICATIONS

* = equal contribution
† = equal advising

Andy Han*, Kristina Fujimoto*, Avidan Shah*, Kiet Nguyen, Kai Xu, Chen Yueh-Han, Ilia Sucholutsky†, Rico Angell†, “On-Policy Consistency Training Improves LLM Safety with Minimal Capability Degradation”, *preprint*.

Avidan Shah, Jannik Brinkmann, Rico Angell, “Mitigating Adaptive Attacks against Reasoning Models with Activation Consistency Training”, *preprint*.

Rico Angell*, Raghav Singhal*, Zachary Horvitz*, Zhou Yu, Rajesh Ranganath, Kathleen McKeown, He He, “Estimating Tail Risks in Language Model Output Distributions”, In *The Forty-Third International Conference on Machine Learning (ICML 2026)*. **Spotlight Presentation.**

Xinpeng Wang*, Nitish Joshi*, Barbara Plank, Rico Angell†, and He He†, “Is It Thinking or Cheating? Detecting Implicit Reward Hacking by Measuring Reasoning Effort”, In *The Fourteenth International Conference on Learning Representations (ICLR 2026)*. **Oral Presentation.**

Rico Angell, Jannik Brinkmann, and He He, “Jailbreak Transferability Emerges from Shared Representations”, In *The Fourteenth International Conference on Learning Representations (ICLR 2026)*.

Chen Yueh-Han, Nitish Joshi, Yulin Chen, Maksym Andriushchenko, Rico Angell[†], and He He[†], “Monitoring Decomposition Attacks in LLMs with Lightweight Sequential Monitors”, In *The Fourteenth International Conference on Learning Representations (ICLR 2026)*.

Adam Karvonen, Benjamin Wright, Can Rager, Rico Angell, Jannik Brinkmann, Logan Smith, Claudio Mayrink Verdun, David Bau, Samuel Marks, “Measuring Progress in Dictionary Learning for Language Model Interpretability with Board Game Models”, In *The 38nd Conference on Neural Information Processing Systems (NeurIPS 2024)* and *ICML 2024 Mechanistic Interpretability Workshop*. **Oral Presentation**.

Rico Angell and Andrew McCallum, “Fast, Scalable, Warm-Start Semidefinite Programming with Spectral Bundling and Sketching”, In *The 41st International Conference on Machine Learning (ICML 2024)*.

Brittany Johnson, Jesse Bartola, Rico Angell, Katherine Keith, Sam Witty, Stephen J Giguere, and Yuriy Brun, “Fairkit, Fairkit, on the Wall, Who’s the Fairest of Them All? Supporting Data Scientists in Training Fair Models”, In *EURO Journal on Decision Processes, 2023*.

Rico Angell, Nicholas Monath, Nishant Yadav, and Andrew McCallum, “Interactive Correlation Clustering with Existential Cluster Constraints”, In *The 39th International Conference on Machine Learning (ICML 2022)*.

Dhruv Agarwal, Rico Angell, Nicholas Monath, and Andrew McCallum, “Entity Linking via Explicit Mention-Mention Coreference Modeling”, In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL 2022)*.

Nishant Yadav, Nicholas Monath, Rico Angell, and Andrew McCallum, “Event and Entity Coreference using Trees to Encode Uncertainty in Joint Decisions”, In *Proceedings of the Fourth Workshop on Computational Models of Reference, Anaphora and Coreference (EMNLP/CRAC 2021)*. **Best Paper Award**.

Rico Angell, Nicholas Monath, Sunil Mohan, Nishant Yadav and Andrew McCallum, “Clustering-based Inference for Biomedical Entity Linking”, In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL 2021)*.

Sunil Mohan, Rico Angell, Nick Monath, and Andrew McCallum, “Low Resource Recognition and Linking of Biomedical Concepts from a Large Ontology”, *Proceedings of the 12th ACM Conference on Bioinformatics, Computational Biology, and Health Informatics (BCB 2021)*

Arthur Feeney*, Rishabh Gupta*, Veronika Thost, Rico Angell, Gayathri Chandu, Yash Adhikari, and Tengfei Ma, “Relation-Dependent Sampling for Multi-Relational Link Prediction”, *ICML 2020 Workshop on Graph Representation Learning and Beyond (ICML/GRL+ 2020)*.

Rico Angell and Daniel Sheldon, “Inferring Latent Velocities from Weather Radar Data using Gaussian Processes”, In *The 32nd Conference on Neural Information Processing Systems (NeurIPS 2018)*.

Rico Angell, Brittany Johnson, Yuriy Brun, and Alexandra Meliou, “Themis: Automatically Testing Software for Discrimination”, In *Proceedings of the 26th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering (ESEC/FSE 2018)*.

Rico Angell and Grant Schoenebeck, “Don’t Be Greedy: Leveraging Community Structure to Find

High Quality Seed Sets for Influence Maximization”, In *The 13th International Conference on Web and Internet Economics (WINE 2017)*.

Rico Angell, Ben Oztalay, and Andrew DeOrio, “A Topological Approach to Hardware Bug Triage”, In *16th International Workshop on Microprocessor and SOC Test and Verification (MTV 2015)*.

HONORS AND AWARDS

New York University

UK AI Security Institute grant via the Alignment Project, Award Value: \$330,476, 2025

University of Massachusetts

NeurIPS 2024 Top Reviewer

Best Workshop Paper Honorable Mention, ICML 2024 Mechanistic Interpretability Workshop

Best Workshop Paper Award, EMNLP/CRAC 2021

Passed Ph.D. candidacy with distinction, 2019

NSF Graduate Research Fellowship, Award Value: \$102,000, 2019

NEAGEP Fellowship (funded by NSF), Award Value: \$25,000, 2017

University of Michigan

Magna Cum Laude

Engineering Honors Program

Summer Undergraduate Research in Engineering Program, Award Value: \$4,200, 2015

Intel Semiconductor Research Corporation Undergraduate Fellowship, Award Value: \$4,000, 2014

University Honors and Engineering Dean’s List: All terms

TALKS

Vector Institute, *Jailbreak Transferability Emerges from Shared Representations*, March 2026

MILA, *Jailbreak Transferability Emerges from Shared Representations*, March 2026

Brown University, *Jailbreak Transferability Emerges from Shared Representations*, February 2026

Northeastern, *Jailbreak Transferability Emerges from Shared Representations*, February 2026

OpenAI, *Jailbreak Transferability Emerges from Shared Representations*, January 2026

UK AISI, *Jailbreak Transferability Emerges from Shared Representations*, January 2026

New York University, *Measuring Progress in Dictionary Learning for Language Model Interpretability with Board Game Models*, July 2024

PROFESSIONAL SERVICE

Reviewer

NeurIPS, ICML, ICLR, ACL ARR, TMLR

TEACHING EXPERIENCE

University of Massachusetts

Teaching Assistant

COMPSCI 696DS, Industry Mentorship Program, Spring 2021 and Lead TA Spring 2022.

University of Michigan

Teaching Assistant

EECS 280, Programming and Introductory Data Structures, Winter 2015.